

B-Tree


Motivation

- So far, we have assumed that we can store an entire data structure in the main memory (**RAM**) of a computer.
- When **data is too large to fit in main memory**, the data structure must reside on disk and the number of disk accesses becomes important.
- A **disk access is expensive** compared to a typical computer instruction.
- Although disks are cheaper and have higher capacity than main memory, they are much, much **slower** because they have moving **mechanical** parts compared to the purely electronic media.
- **One** disk access is worth **thousands** of instructions.
- The number of disk accesses will dominate the running time.

Motivation Cont..

- Secondary memory (disk) is divided into equal-sized **blocks** (typical sizes are 512, 2048, 4096 or 8192 bytes)
- The basic **I/O** operation transfers the contents of one disk block to/from main memory.
- Our goal is to devise a multiway search tree that will **minimize** file accesses.

m-ary Trees




- m-ary search tree allows **m-way branching (m children)**.
- A **node** in a m-tree contains **multiple keys** (K1,K2,K3 etc).
- Each piece of data stored is called a "**key**", because each key is unique and can occur in only one location.
- The **keys** in a node serve as **dividing points** separating the range of keys.
- In applications there would be a record of data associated with each key (e.g, student record, phone number).
- Order of subtrees is based on parent node keys.
- If each node has **m children** and there are **n keys** then the average time taken to search, insert and delete from the tree is **$\log_m n$** .

B-Tree


- Ideally, a tree will be balanced and the height will be $\log n$ where n is the number of nodes in the tree. To ensure that the height of the tree is as small as possible and therefore provide the best running time, a **balanced tree** structure like a AVL tree, b-tree, or red-black tree must be used.
- **B-tree** is a self-balancing tree data structure that keeps data **sorted** and allows **searches**, sequential access, insertions, and deletions in logarithmic time.
- The B-tree is a **generalization of a binary search tree** in that a node can have more than two children.
- The idea is that we leave some **key spaces open**. So an insert of a new key is done using available space (most cases).
- Unlike self-balancing binary search trees, the B-tree is optimized for systems that read and write **large blocks of data**.

- The B-tree algorithm **minimizes** the number of times a medium must be accessed to locate a desired record, thereby **speeding** up the process.
- The B-tree algorithms **copy selected pages from disk** into main memory as needed and write back onto disk the pages that have changed. B-tree algorithms keep only a constant number of pages in main memory at any time; thus, the size of main memory does not limit the size of B-trees that can be handled.
- In a practical B-tree, there can be thousands, millions, or **billions** of records.
- Real-world B-trees are of higher order (32, 64, 128, or more). For a large B-tree stored on a disk, we often see **branching factors** between 50 and 2000.

- A B-tree **node** is usually as large as a whole **disk page**, and this size limits the number of children a B-tree node can have.
- B-trees are a good example of a data structure for external memory. It is commonly used in **databases** and **filesystems**. (e.g., Used in Mac, NTFS, OS2 for file structure. SQL Server [https://technet.microsoft.com/en-us/library/ms177443\(v=sql.105\).aspx](https://technet.microsoft.com/en-us/library/ms177443(v=sql.105).aspx)).
- Real world database indexes with millions of records have a **tree depth of four or five**.




Binary tree of 31
nodes has 5 levels



A 5-ary tree of
31 nodes has
only 3 levels

Binary Tree vs B-Tree

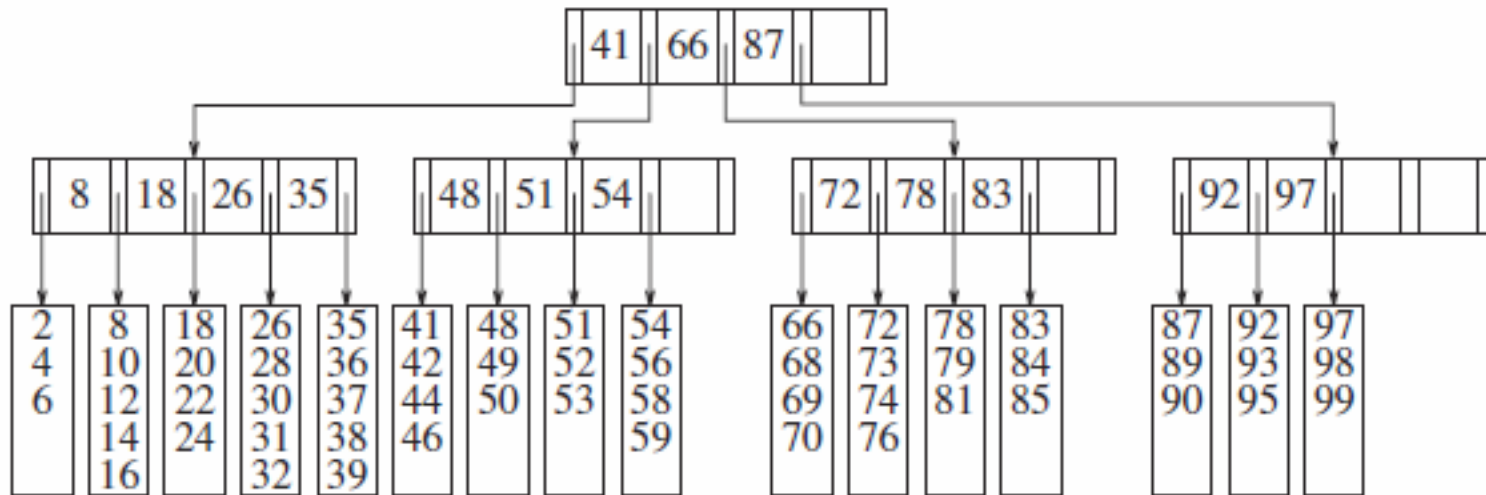


- **B-trees save time** by using nodes with **many branches**, compared with **binary trees**, in which each node has only two children. When there are many children per node, a record can be found by passing through fewer nodes than if there are two children per node.
- The tradeoff is that the **decision process** at each node is more complicated in a B-tree as compared with a binary tree. A program is required to execute the operations in a B-tree. But this program is stored in **RAM**, so it runs fast.

Properties of a B-Tree


- **Order** of a tree = maximum number of **children** per node.
- B-tree of **order m** has the following properties:
 1. The root is either a leaf or has between 2 and m children.
 2. Every node has at most **m children (pointers)**.
 3. Every node except for the root and the leaves has **at least $\lceil m/2 \rceil$ children**.
 4. A non-leaf node stores up to **m-1 keys**.
 5. All leaves appear at the same level.

A B-tree of order 5



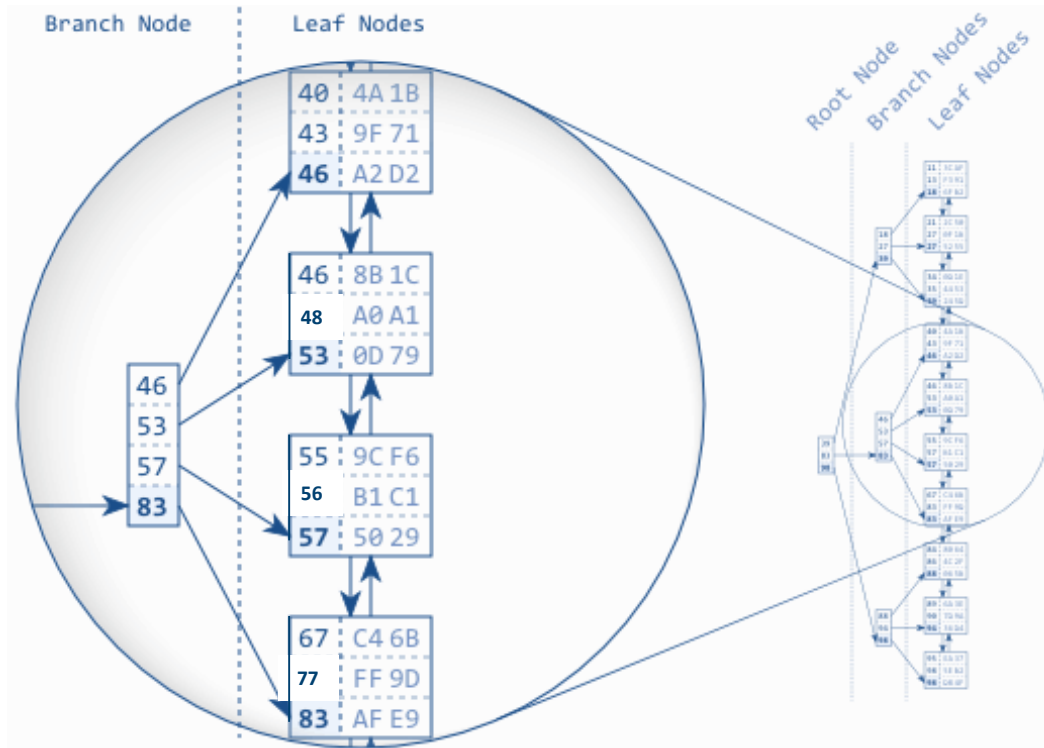
Notice that all non-leaf nodes have between 3 and 5 children (and thus between 2 and 4 keys).

How many keys a B-tree can store with a branching factor of 1001 and height 2?



- Max number of keys at depth d in a B-tree = $m^d * k$
 - m = order of tree
 - k = max number of keys per node
 - d = depth
- Since we can keep the **root node permanently in main memory**, we can find any key in this tree by making at most only two disk accesses.


B-tree Structure



- An index with 30 entries.
- An implementation of B-Tree that is used in some databases: Each branch node entry corresponds to the biggest value in the respective leaf node.

B-Tree Traversal

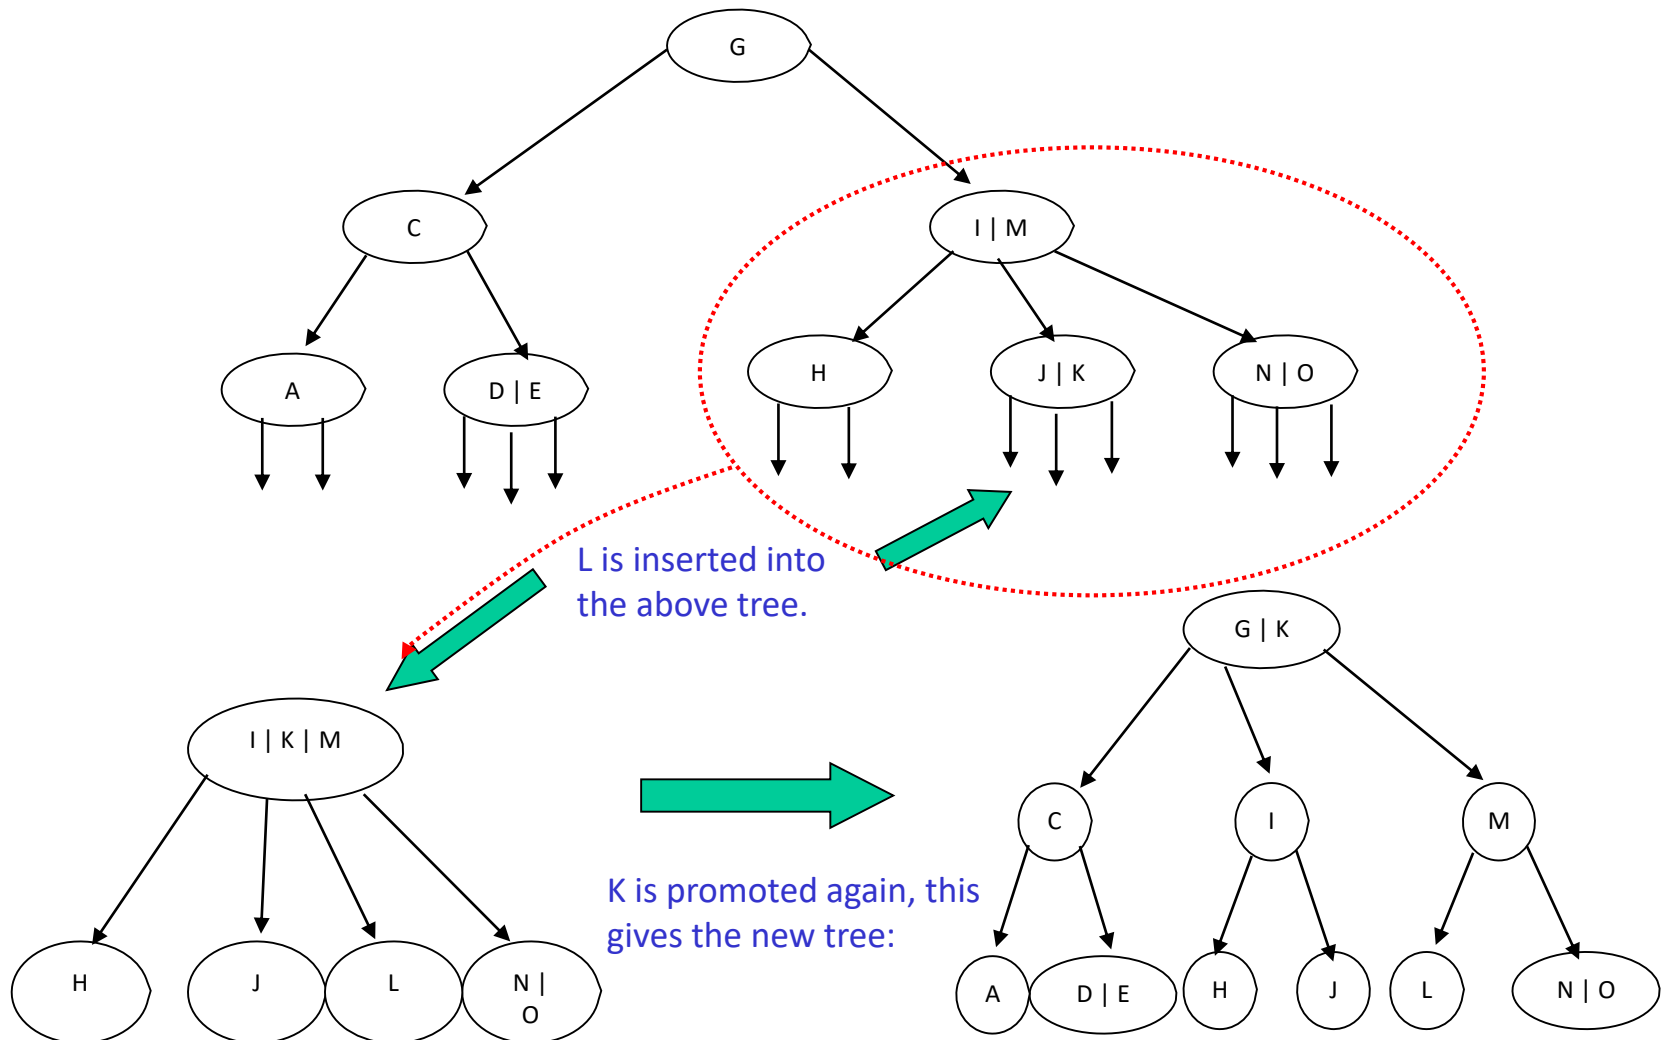
The figure shows an index fragment to illustrate a search for the key “56”. The tree traversal starts at the root node on the left-hand side. Each entry is processed in ascending order **until a value is greater than or equal to (\geq) the search term (56)**. In the figure it is the entry 83. The database follows the reference to the corresponding branch node and repeats the procedure until the tree traversal reaches a leaf node.




Insertion

- Insert k into B-tree of order m .
 - We find the insertion point (in a leaf) by doing a **search**.
 - **If there is room then enter k** , keeping the node's elements ordered.
 - Else, **promote the middle key** to the parent & **split** the node into 2 nodes around the middle key.
 - If the splitting backs up to the root, then make a new root containing the middle key.
- Note: the tree grows from the leaves, **balance is always maintained**.


Insertion Example – B-Tree of order 3




Splitting Nodes





- Middle key is promoted
- Join its parent or create a new root



Insert the numbers 1 to 7 into a B-Tree of order 3



Show the final B-Tree after inserting keys B, Q, L, and F into the following B-tree of order 6




Insert the following letters into an empty B-tree of order 5:
C N G A H E K Q M F W L T Z D P R X Y S


- Order 5 means that a node can have a maximum of 5 children and 4 keys.
- All nodes other than the root must have a minimum of 2 keys.
- The first 4 letters get inserted into the same node, resulting in this picture:

A	C	G	N


When we try to insert the H, we find no room in this node, so we split it into 2 nodes, moving the median item G up into a new root node.




Inserting E, K, and Q proceeds without requiring any splits:




Inserting M requires a split. Note that M happens to be the median key and so is moved up into the parent node.




The letters F, W, L, and T are then added without needing any split.




When Z is added, the rightmost leaf must be split. The median item T is moved up into the parent node. Note that by moving up the median key, the tree is kept fairly balanced, with 2 keys in each of the resulting nodes.




The insertion of D causes the leftmost leaf to be split. D happens to be the median key and so is the one moved up into the parent node. The letters P, R, X, and Y are then added without any need of splitting:



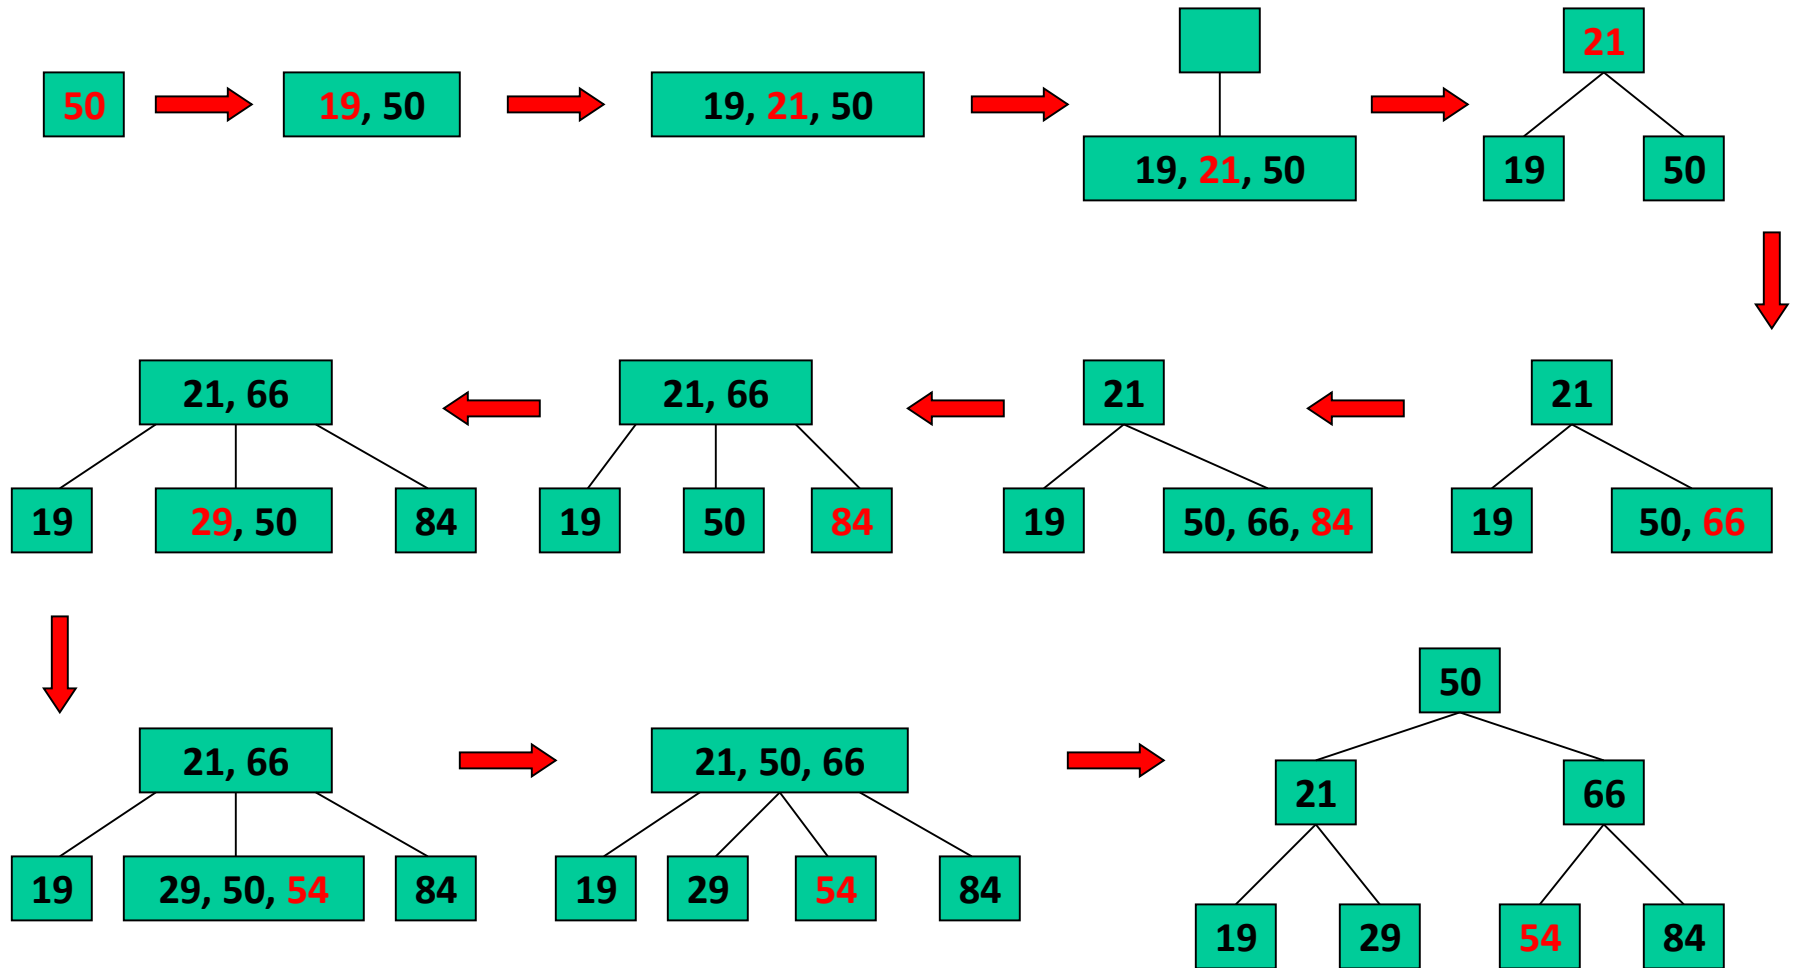
Finally, when S is added, the node with N, P, Q, and R splits, sending the median Q up to the parent. However, the parent node is full, so it splits, sending the median M up to form a new root node. Note how the 3 pointers from the old parent node stay in the revised node that contains D and G.



Insert the following letters into an empty B-tree of order 5:
A G F B K D H M J E S I R X C L N T U P



Insert 50, 19, 21, 66, 84, 29, and 54 into a B-Tree with order 3



Useful links

- B-Tree applet
<https://www.cs.usfca.edu/~galles/visualization/BTree.html>
- B-Tree talk
- <http://www.csanimated.com/animation.php?t=B-tree>